

airline-data-eda-and-prediction

June 27, 2024

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
from matplotlib import pyplot as plt
import networkx as nx
```

```
[2]: df = pd.read_csv('/kaggle/input/flight-price-data/flight_dataset.csv')
```

1 Data Cleaning

```
[3]: df.head()
```

```
[3]:
```

	Airline	Source	Destination	Total_Stops	Price	Date	Month	Year	\
0	IndiGo	Banglore	New Delhi	0	3897	24	3	2019	
1	Air India	Kolkata	Banglore	2	7662	1	5	2019	
2	Jet Airways	Delhi	Cochin	2	13882	9	6	2019	
3	IndiGo	Kolkata	Banglore	1	6218	12	5	2019	
4	IndiGo	Banglore	New Delhi	1	13302	1	3	2019	

	Dep_hours	Dep_min	Arrival_hours	Arrival_min	Duration_hours	\
0	22	20	1	10	2	
1	5	50	13	15	7	
2	9	25	4	25	19	
3	18	5	23	30	5	
4	16	50	21	35	4	

	Duration_min
0	50
1	25
2	0
3	25
4	45

```
[4]: df.describe()
```

```
[4]:
```

	Total_Stops	Price	Date	Month	Year \
count	10683.000000	10683.000000	10683.000000	10683.000000	10683.0
mean	0.824207	9087.064121	13.508378	4.708602	2019.0
std	0.675199	4611.359167	8.479277	1.164357	0.0
min	0.000000	1759.000000	1.000000	3.000000	2019.0
25%	0.000000	5277.000000	6.000000	3.000000	2019.0
50%	1.000000	8372.000000	12.000000	5.000000	2019.0
75%	1.000000	12373.000000	21.000000	6.000000	2019.0
max	4.000000	79512.000000	27.000000	6.000000	2019.0

	Dep_hours	Dep_min	Arrival_hours	Arrival_min \
count	10683.000000	10683.000000	10683.000000	10683.000000
mean	12.490686	24.411214	13.348778	24.690630
std	5.748650	18.767980	6.859125	16.506036
min	0.000000	0.000000	0.000000	0.000000
25%	8.000000	5.000000	8.000000	10.000000
50%	11.000000	25.000000	14.000000	25.000000
75%	18.000000	40.000000	19.000000	35.000000
max	23.000000	55.000000	23.000000	55.000000

	Duration_hours	Duration_min
count	10683.000000	10683.000000
mean	10.246560	28.327249
std	8.494988	16.946113
min	1.000000	0.000000
25%	2.000000	15.000000
50%	8.000000	30.000000
75%	15.000000	45.000000
max	47.000000	55.000000

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10683 entries, 0 to 10682
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Airline         10683 non-null  object
1   Source          10683 non-null  object
2   Destination     10683 non-null  object
3   Total_Stops     10683 non-null  int64
4   Price           10683 non-null  int64
5   Date            10683 non-null  int64
6   Month           10683 non-null  int64
7   Year            10683 non-null  int64
8   Dep_hours       10683 non-null  int64
9   Dep_min         10683 non-null  int64
```

```

10  Arrival_hours    10683 non-null   int64
11  Arrival_min      10683 non-null   int64
12  Duration_hours   10683 non-null   int64
13  Duration_min     10683 non-null   int64
dtypes: int64(11), object(3)
memory usage: 1.1+ MB

```

```
[6]: df.isna().sum()
```

```

[6]: Airline          0
     Source           0
     Destination      0
     Total_Stops       0
     Price            0
     Date             0
     Month            0
     Year             0
     Dep_hours        0
     Dep_min          0
     Arrival_hours    0
     Arrival_min      0
     Duration_hours   0
     Duration_min     0
     dtype: int64

```

```
[7]: df.duplicated().sum()
```

```
[7]: 222
```

```

[8]: # Dropping duplicated rows
     df.drop_duplicates(inplace = True)

```

```
[9]: df.shape
```

```
[9]: (10461, 14)
```

```
[10]: df.head()
```

```

[10]:
   Airline  Source Destination  Total_Stops  Price  Date  Month  Year  \
0   IndiGo  Bangalore  New Delhi           0   3897   24     3  2019
1  Air India  Kolkata   Bangalore           2   7662    1     5  2019
2  Jet Airways    Delhi    Cochin           2  13882    9     6  2019
3   IndiGo  Kolkata   Bangalore           1   6218   12     5  2019
4   IndiGo  Bangalore  New Delhi           1  13302    1     3  2019


   Dep_hours  Dep_min  Arrival_hours  Arrival_min  Duration_hours  \
0          22      20              1           10              2

```

1	5	50	13	15	7
2	9	25	4	25	19
3	18	5	23	30	5
4	16	50	21	35	4

	Duration_min
0	50
1	25
2	0
3	25
4	45

```
[11]: plt.figure(figsize=(12,8))
sns.boxplot(df)
plt.show()
```



```
[12]: # Removing the outliers from price column
```

```
q1 = df['Price'].quantile(0.25)
q3 = df['Price'].quantile(0.75)

iqr = q3 - q1

ll = q1 - 1.5 * iqr
```

```
ul = q3 + 1.5 * iqr

outliers = (df['Price'] < ll ) |( df['Price'] > ul)
df = df[~outliers]
```

[13]: *# Removing the outliers from Duration_hours column*

```
q1 = df['Duration_hours'].quantile(0.25)
q3 = df['Duration_hours'].quantile(0.75)

iqr = q3 - q1

ll = q1 - 1.5 * iqr
ul = q3 + 1.5 * iqr

outliers = (df['Duration_hours'] < ll ) |( df['Duration_hours'] > ul)
df = df[~outliers]
```

[14]: *# Removing the outliers from Total_Stops column*


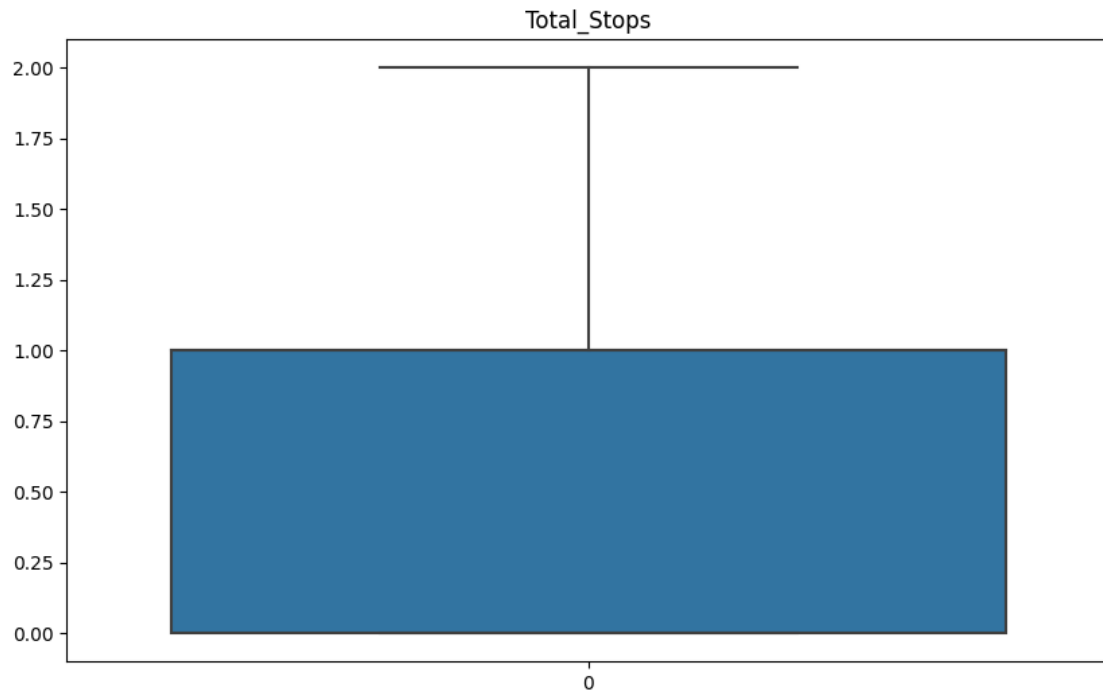
```
q1 = df['Total_Stops'].quantile(0.25)
q3 = df['Total_Stops'].quantile(0.75)



iqr = q3 - q1

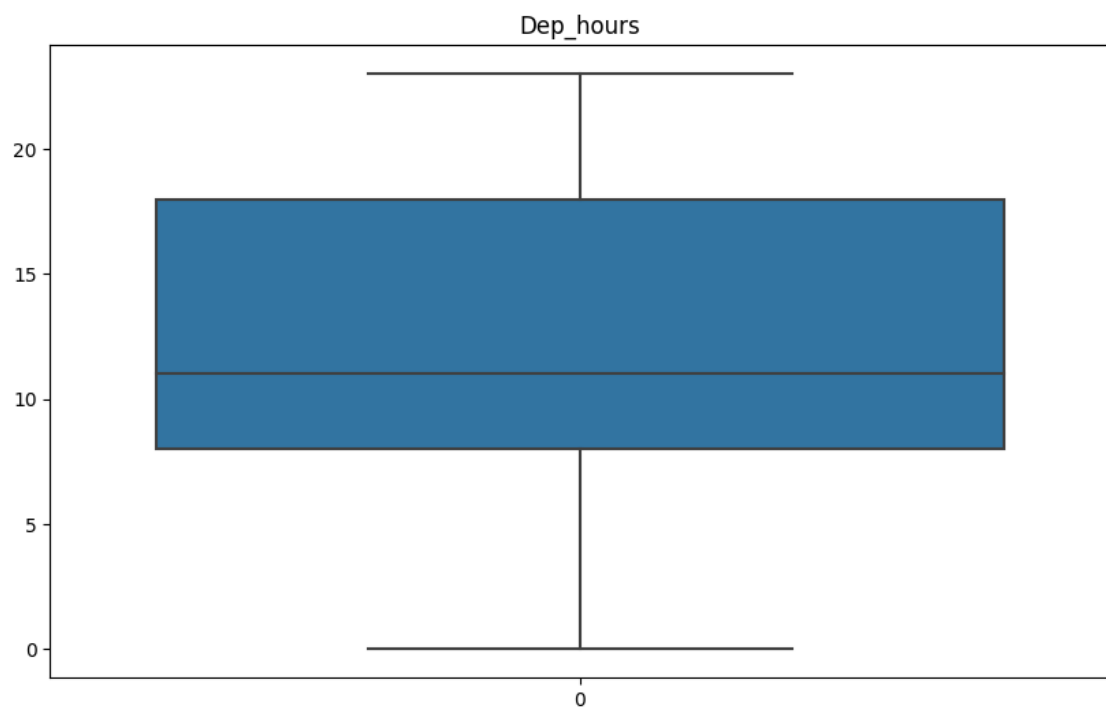

ll = q1 - 1.5 * iqr
ul = q3 + 1.5 * iqr



outliers = (df['Total_Stops'] < ll ) |( df['Total_Stops'] > ul)
df = df[~outliers]
```

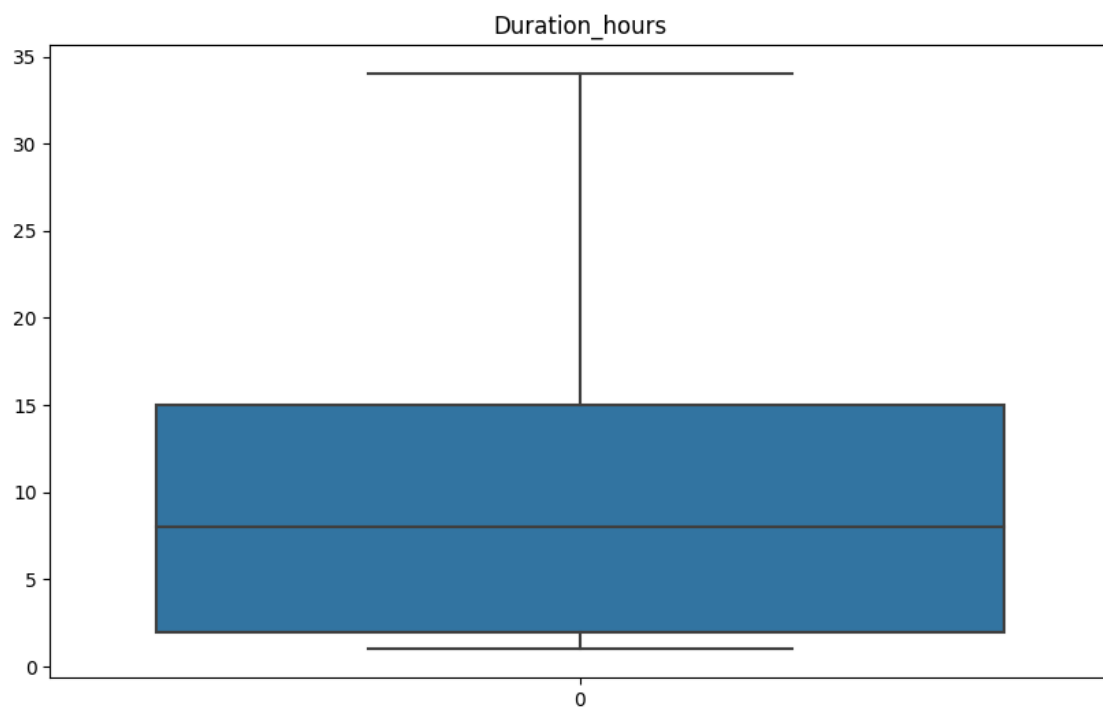
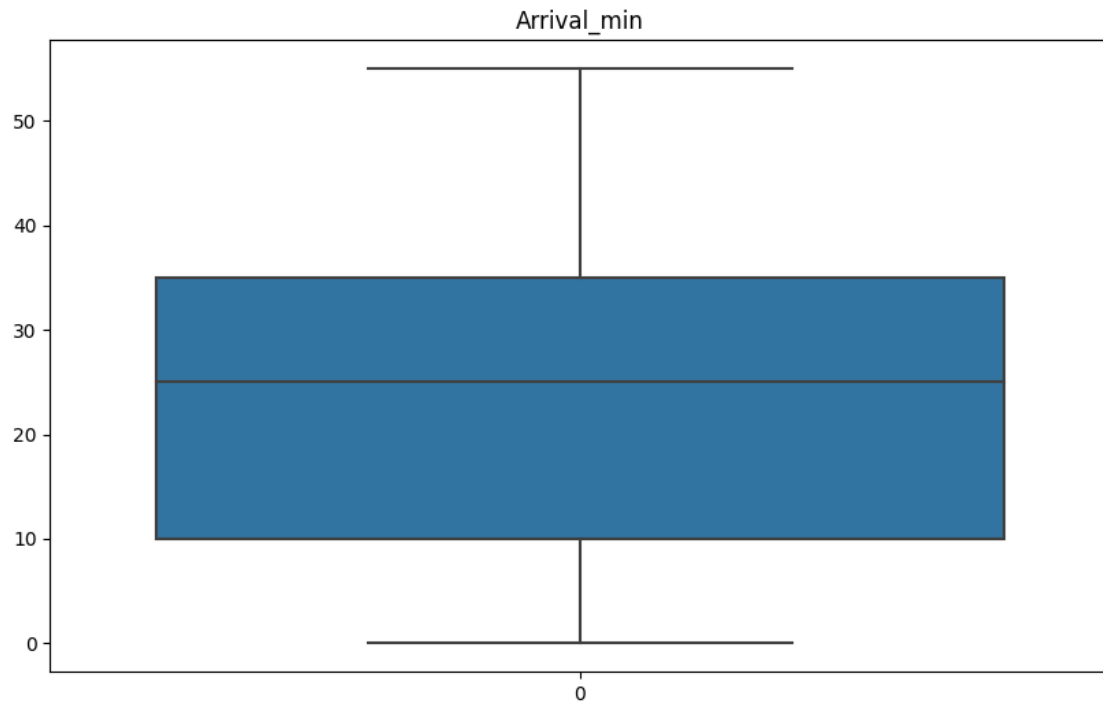
[15]: `num_cols = df.select_dtypes(include=['int64'])`
`for i in num_cols.columns:`
 `plt.figure(figsize=(10,6))`
 `sns.boxplot(df[i])`
 `plt.title(i)`
 `plt.show()`

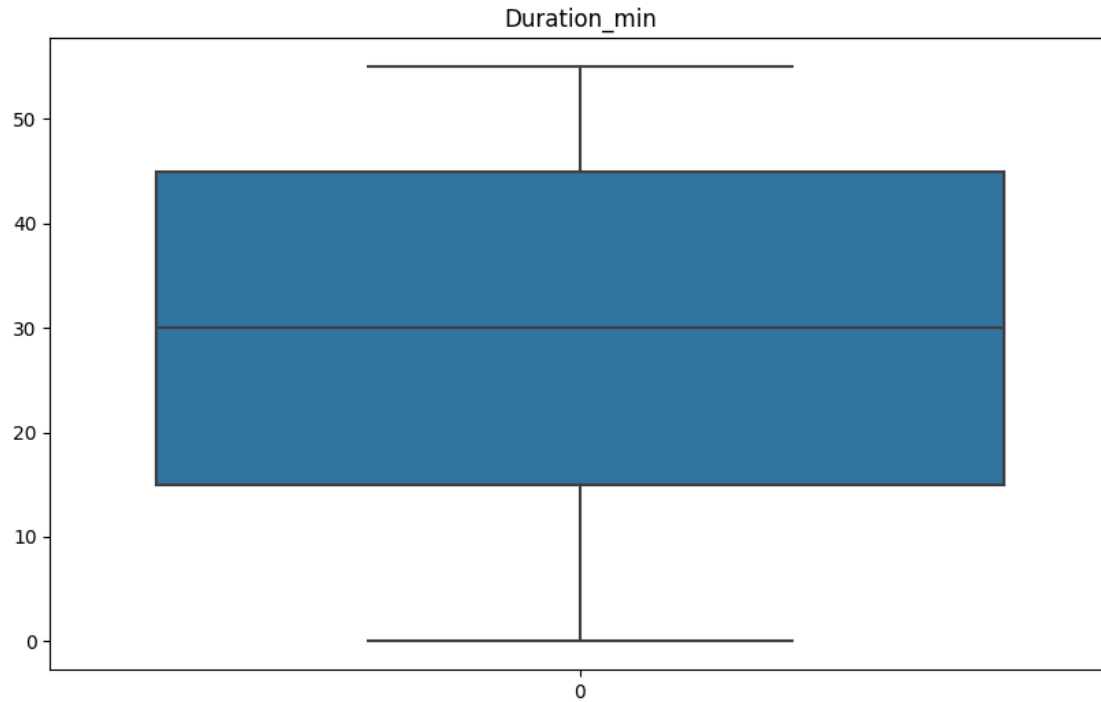














2 EDA

```
[16]: # histogram for price column
plt.figure(figsize=(12,10))
sns.histplot(df['Price'], kde = True , color = 'orange')
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
with pd.option_context('mode.use_inf_as_na', True):



```
[17]: df['Airline'].value_counts().plot(kind = 'bar')  
plt.show()
```



```
[18]: # Dropping the Year column as it has only one value(2019)
df.drop(columns = ['Year'], inplace = True)
```

```
[19]: df.head()
```

```
[19]:
```

	Airline	Source	Destination	Total_Stops	Price	Date	Month	\
0	IndiGo	Banglore	New Delhi	0	3897	24	3	
1	Air India	Kolkata	Banglore	2	7662	1	5	
2	Jet Airways	Delhi	Cochin	2	13882	9	6	

3	IndiGo	Kolkata	Banglore	1	6218	12	5
4	IndiGo	Banglore	New Delhi	1	13302	1	3

	Dep_hours	Dep_min	Arrival_hours	Arrival_min	Duration_hours	\
0	22	20	1	10		2
1	5	50	13	15		7
2	9	25	4	25		19
3	18	5	23	30		5
4	16	50	21	35		4

	Duration_min
0	50
1	25
2	0
3	25
4	45

```
[20]: # Checking the most popular routes
route_freq = df.groupby(['Source', 'Destination']).size().
        ↪reset_index(name='count')
route_freq = route_freq.sort_values(by = 'count' , ascending = False)
```

```
[21]: route_freq.head()
```

	Source	Destination	count
3	Delhi	Cochin	4265
4	Kolkata	Banglore	2846
0	Banglore	Delhi	1265
1	Banglore	New Delhi	826
5	Mumbai	Hyderabad	688

```
[22]: G = nx.from_pandas_edgelist(route_freq, 'Source', 'Destination', ['count'],
        ↪create_using=nx.DiGraph())
plt.figure(figsize=(15, 10))

pos = nx.spring_layout(G, k=0.5, iterations=50)
nx.draw_networkx_nodes(G, pos, node_size=3000, node_color='skyblue',
        ↪edgecolors='black')


nx.draw_networkx_edges(G, pos, arrowstyle='->', arrowsize=20,
        ↪edge_color='gray', width=2)

nx.draw_networkx_labels(G, pos, font_size=12, font_weight='bold')

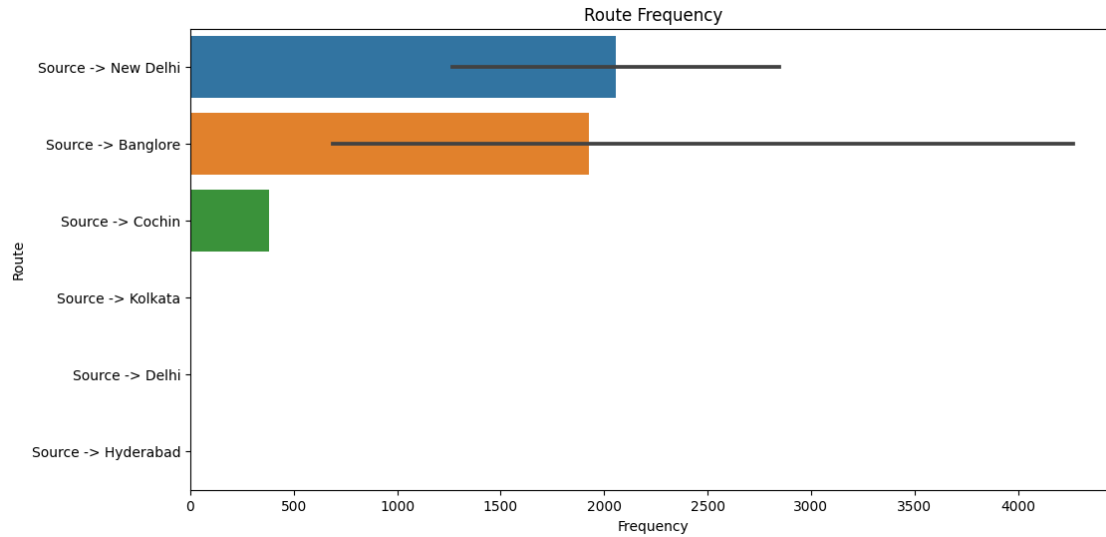
edge_labels = nx.get_edge_attributes(G, 'count')
nx.draw_networkx_edge_labels(G, pos, edge_labels=edge_labels, font_size=10,
        ↪font_color='red')
```

```
plt.title('Airline Route Network', size=20)

plt.axis('off')
plt.show()
```



```
[23]: plt.figure(figsize=(12, 6))
sns.barplot(data=route_freq, x='count', y='Source' + ' -> ' + df['Destination'])
plt.xlabel('Frequency')
plt.ylabel('Route')
plt.title('Route Frequency')
plt.show()
```



```
[24]: df.head()
```


```
[24]:
```

	Airline	Source	Destination	Total_Stops	Price	Date	Month	\
0	IndiGo	Banglore	New Delhi	0	3897	24	3	
1	Air India	Kolkata	Banglore	2	7662	1	5	
2	Jet Airways	Delhi	Cochin	2	13882	9	6	
3	IndiGo	Kolkata	Banglore	1	6218	12	5	
4	IndiGo	Banglore	New Delhi	1	13302	1	3	

	Dep_hours	Dep_min	Arrival_hours	Arrival_min	Duration_hours	\
0	22	20	1	10	2	
1	5	50	13	15	7	
2	9	25	4	25	19	
3	18	5	23	30	5	
4	16	50	21	35	4	

	Duration_min
0	50
1	25
2	0
3	25
4	45

```
[25]: # Total Stops vs airline
df.groupby('Airline').Total_Stops.mean().sort_values(ascending = False).
    plot(kind = 'bar', color = 'orange')
plt.show()
```

```
[26]: # Total Stops vs source and destination
route_stops = df.groupby(['Source', 'Destination']).Total_Stops.mean().
    ↪reset_index().sort_values(by='Total_Stops', ascending=False)
```

```
[27]: route_stops
```

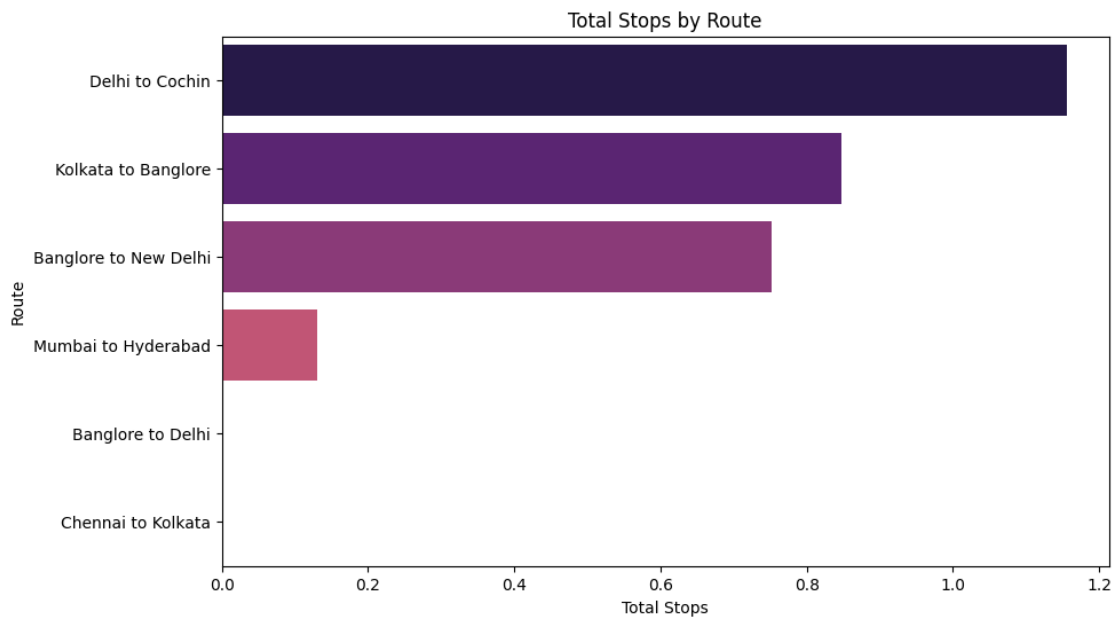
```
[27]:   Source Destination  Total_Stops
3    Delhi      Cochin    1.155451
```

4	Kolkata	Banglore	0.847505
1	Banglore	New Delhi	0.751816
5	Mumbai	Hyderabad	0.130814
0	Banglore	Delhi	0.000000
2	Chennai	Kolkata	0.000000

```
[28]: plt.figure(figsize=(10, 6))
sns.barplot(x='Total_Stops', y=route_stops['Source'] + ' to ' +
            route_stops['Destination'], data=route_stops, palette='magma')

plt.xlabel('Total Stops')
plt.ylabel('Route')
plt.title('Total Stops by Route')

plt.show()
```



```
[29]: df.head()
```

```
[29]:
```

	Airline	Source	Destination	Total_Stops	Price	Date	Month	\
0	IndiGo	Banglore	New Delhi	0	3897	24	3	
1	Air India	Kolkata	Banglore	2	7662	1	5	
2	Jet Airways	Delhi	Cochin	2	13882	9	6	
3	IndiGo	Kolkata	Banglore	1	6218	12	5	
4	IndiGo	Banglore	New Delhi	1	13302	1	3	

	Dep_hours	Dep_min	Arrival_hours	Arrival_min	Duration_hours	\
--	-----------	---------	---------------	-------------	----------------	---


0	22	20	1	10	2
1	5	50	13	15	7
2	9	25	4	25	19
3	18	5	23	30	5
4	16	50	21	35	4

	Duration_min
0	50
1	25
2	0
3	25
4	45

```
[30]: # Airline vs Price
airline_price = df.groupby('Airline').Price.mean().sort_values(ascending =
↪False)
```

```
[31]: plt.figure(figsize=(10,6))
sns.barplot(x = airline_price.index , y = airline_price.values ,
↪palette='magma')

plt.xticks(rotation = 90)
plt.show()
```



```
[32]: df.head()
```

```
[32]:
```

	Airline	Source	Destination	Total_Stops	Price	Date	Month	\
0	IndiGo	Banglore	New Delhi	0	3897	24	3	
1	Air India	Kolkata	Banglore	2	7662	1	5	
2	Jet Airways	Delhi	Cochin	2	13882	9	6	
3	IndiGo	Kolkata	Banglore	1	6218	12	5	
4	IndiGo	Banglore	New Delhi	1	13302	1	3	

	Dep_hours	Dep_min	Arrival_hours	Arrival_min	Duration_hours	\
0	22	20	1	10	2	
1	5	50	13	15	7	
2	9	25	4	25	19	
3	18	5	23	30	5	
4	16	50	21	35	4	

	Duration_min
0	50
1	25
2	0
3	25
4	45

```
[33]: # Routes vs Price
route_price = df.groupby(['Source', 'Destination']).Price.mean().
    ↪reset_index(name='price')
route_price = route_price.sort_values(by = 'price' , ascending = False)
```


```
[34]: route_price
```

```
[34]:
```

	Source	Destination	price
1	Banglore	New Delhi	10374.016949
3	Delhi	Cochin	10350.422509
4	Kolkata	Banglore	9116.026001
0	Banglore	Delhi	5143.918577
5	Mumbai	Hyderabad	4865.125000
2	Chennai	Kolkata	4789.892388

```
[35]: plt.figure(figsize=(10,6))
sns.barplot(x = route_price['Source'] +" to "+ route_price['Destination'], y =
    ↪route_price['price'] , palette='pastel')

plt.xticks(rotation = 45)
plt.show()
```



```
[36]: # combining the Duration_hours and Duration_min column and converting it into
      ↪ minutes for easier analysis
```

```
df['duration'] = (df['Duration_hours']*60) + df['Duration_min']
```

```
[37]: df.drop(columns=['Duration_hours','Duration_min'], inplace = True)
```

```
[38]: # Routes vs duration
```

```
route_dur = df.groupby(['Source','Destination']).duration.mean().
```

```
      ↪ reset_index(name='duration')
```

```
route_dur = route_dur.sort_values(by = 'duration' , ascending = False)
```


```
[39]: route_dur
```

```
[39]:
```

	Source	Destination	duration
3	Delhi	Cochin	774.990621
4	Kolkata	Banglore	742.127547
1	Banglore	New Delhi	632.929782
5	Mumbai	Hyderabad	175.508721
0	Banglore	Delhi	171.695652
2	Chennai	Kolkata	139.619423

```
[40]: plt.figure(figsize=(10,6))
sns.barplot(y = route_dur['Source'] + " - " + route_dur['Destination'], x =
    ↳route_dur['duration'] , palette='cool')

plt.xticks(rotation = 45)
plt.show()
```




```
[41]: df_sorted = df.sort_values(by='duration')

plt.figure(figsize=(12, 10))
plt.plot(df_sorted['duration'], df_sorted['Price'], marker='o',
    ↳color='skyblue', linewidth=2, markersize=8)

plt.xlabel('Duration')
plt.ylabel('Price')
plt.title('Duration vs Price')

plt.grid(True)
plt.show()
```



```
[42]: df.head()
```

```
[42]:
```

	Airline	Source	Destination	Total_Stops	Price	Date	Month	\
0	IndiGo	Banglore	New Delhi	0	3897	24	3	
1	Air India	Kolkata	Banglore	2	7662	1	5	
2	Jet Airways	Delhi	Cochin	2	13882	9	6	
3	IndiGo	Kolkata	Banglore	1	6218	12	5	
4	IndiGo	Banglore	New Delhi	1	13302	1	3	

	Dep_hours	Dep_min	Arrival_hours	Arrival_min	duration
0	22	20	1	10	170
1	5	50	13	15	445
2	9	25	4	25	1140
3	18	5	23	30	325
4	16	50	21	35	285

```
[43]: # month
df.Month.value_counts()
```


[43]: Month

5	3379
6	3292
3	2525
4	1075

Name: count, dtype: int64

```
[44]: df.Date.value_counts().sort_values(ascending = False)
```

[44]: Date

9	1359
6	1250
27	1081
21	1075
24	1013
1	976
15	960
12	934
3	814
18	809

Name: count, dtype: int64

```
[45]: df.Dep_hours.value_counts().sort_values(ascending = False)
```

[45]: Dep_hours

9	877
7	850
17	684
8	679
6	666
20	645
11	561
19	541
10	522
5	520
14	499
21	486
16	452
18	437
13	411
22	361
15	317
2	194
12	171
4	168
23	131
0	38

```
1      37
3      24
Name: count, dtype: int64
```

```
[46]: sns.set(style="whitegrid", palette="pastel")

# Customize the pairplot
pairplot = sns.pairplot(df,
                        kind="scatter",
                        diag_kind="kde",
                        markers="o",
                        plot_kws={'alpha':0.6, 's':80, 'edgecolor':'k'},
                        diag_kws={'fill': True})

# Add titles and adjust the plot
pairplot.fig.suptitle("Pairplot of Airline Data", y=1.02) # y=1.02 to move the
↳title a bit above the plot
pairplot.fig.set_size_inches(12, 10)

# Show the plot
plt.show()
```

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```

```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```


```
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
```

use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.

```
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```


```
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):
```



```
[47]: # heatmap for the dataset
plt.figure(figsize=(10, 8))
num_cols = df.select_dtypes('int64')
heatmap = sns.heatmap(num_cols.corr(), annot=True, cmap='coolwarm', fmt=".2f",
↳linewidths=.5)
```

```
heatmap.set_title('Correlation Heatmap')
plt.show()
```



```
[48]: # price column is highly co related to duration and total stops and that makes sense
```

3 Model building and evaluation

```
[49]: df.head()
```

```
[49]:
```

	Airline	Source	Destination	Total_Stops	Price	Date	Month	\
0	IndiGo	Banglore	New Delhi	0	3897	24	3	
1	Air India	Kolkata	Banglore	2	7662	1	5	
2	Jet Airways	Delhi	Cochin	2	13882	9	6	
3	IndiGo	Kolkata	Banglore	1	6218	12	5	

4	IndiGo	Banglore	New Delhi	1	13302	1	3
---	--------	----------	-----------	---	-------	---	---

	Dep_hours	Dep_min	Arrival_hours	Arrival_min	duration
0	22	20	1	10	170
1	5	50	13	15	445
2	9	25	4	25	1140
3	18	5	23	30	325
4	16	50	21	35	285

```
[50]: from sklearn.preprocessing import OneHotEncoder,MinMaxScaler
```

```
[51]: # Normalization

def normalize_columns(df, columns):
    for col in columns:
        # Min-max normalization: (x - min) / (max - min)
        min_val = df[col].min()
        max_val = df[col].max()
        df[col] = (df[col] - min_val) / (max_val - min_val)

columns_to_normalize = ['Total_Stops', 'Date',
                        'Month', 'Dep_hours', 'Dep_min', 'Arrival_hours', 'Arrival_min',
                        'duration']

normalize_columns(df, columns_to_normalize)
```

```
[52]: # One-Hot-Encoding of categorical columns
categorical_cols = df.select_dtypes(include=['object']).columns

df = pd.get_dummies(df, columns=categorical_cols, drop_first=True)
```

```
[53]: df.head()
```

```
[53]:
```

	Total_Stops	Price	Date	Month	Dep_hours	Dep_min	Arrival_hours	\
0	0.0	3897	0.884615	0.000000	0.956522	0.363636	0.043478	
1	1.0	7662	0.000000	0.666667	0.217391	0.909091	0.565217	
2	1.0	13882	0.307692	1.000000	0.391304	0.454545	0.173913	
3	0.5	6218	0.423077	0.666667	0.782609	0.090909	1.000000	
4	0.5	13302	0.000000	0.000000	0.695652	0.909091	0.913043	

	Arrival_min	duration	Airline_Air	India	...	\
0	0.181818	0.047619		False	...	
1	0.272727	0.185464		True	...	
2	0.454545	0.533835		False	...	
3	0.545455	0.125313		False	...	
4	0.636364	0.105263		False	...	

	Airline_Vistara	Premium economy	Source_Chennai	Source_Delhi	\
0		False	False	False	
1		False	False	False	
2		False	False	True	
3		False	False	False	
4		False	False	False	

	Source_Kolkata	Source_Mumbai	Destination_Cochin	Destination_Delhi	\
0	False	False	False	False	
1	True	False	False	False	
2	False	False	True	False	
3	True	False	False	False	
4	False	False	False	False	

	Destination_Hyderabad	Destination_Kolkata	Destination_New Delhi
0	False	False	True
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	True

[5 rows x 28 columns]

```
[54]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error,
    median_absolute_error, r2_score, explained_variance_score

X = df.drop(columns=['Price'])
y = df['Price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

models = {
    'Linear Regression': LinearRegression(),
    'Ridge Regression': Ridge(),
    'Lasso Regression': Lasso(),
    'Decision Tree Regressor': DecisionTreeRegressor(random_state=42),
    'Random Forest Regressor': RandomForestRegressor(random_state=42),
    'Gradient Boosting Regressor': GradientBoostingRegressor(random_state=42)
}

for model_name, model in models.items():
```

```

print(f"Training {model_name}...")
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
medae = median_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
evs = explained_variance_score(y_test, y_pred)

print(f"\n{model_name} Results:")
print(f"Mean Squared Error (MSE): {mse:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"Median Absolute Error: {medae:.2f}")
print(f"R^2 Score: {r2:.2f}")
print(f"Explained Variance Score: {evs:.2f}")
print("\n")

plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, color='blue')
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)],
         linestyle='--', color='red')
plt.xlabel('Actual Price')
plt.ylabel('Predicted Price')
plt.title(f'{model_name} - Actual vs Predicted Price')
plt.grid(True)
plt.show()

```

Training Linear Regression...

Linear Regression Results:

Mean Squared Error (MSE): 5684185.75

Mean Absolute Error (MAE): 1793.79

Median Absolute Error: 1313.00

R² Score: 0.64

Explained Variance Score: 0.64



Training Ridge Regression...

Ridge Regression Results:


Mean Squared Error (MSE): 5672244.66

Mean Absolute Error (MAE): 1792.40

Median Absolute Error: 1311.73

R^2 Score: 0.64

Explained Variance Score: 0.64



Training Lasso Regression...

Lasso Regression Results:

Mean Squared Error (MSE): 5668318.33

Mean Absolute Error (MAE): 1791.05

Median Absolute Error: 1302.68


R^2 Score: 0.64

Explained Variance Score: 0.64




Training Decision Tree Regressor...

Decision Tree Regressor Results:
Mean Squared Error (MSE): 4774660.00
Mean Absolute Error (MAE): 1317.87
Median Absolute Error: 525.00
 R^2 Score: 0.70
Explained Variance Score: 0.70



Training Random Forest Regressor...

Random Forest Regressor Results:
Mean Squared Error (MSE): 3122732.64
Mean Absolute Error (MAE): 1148.94
Median Absolute Error: 613.32
 R^2 Score: 0.80
Explained Variance Score: 0.80



Training Gradient Boosting Regressor...

Gradient Boosting Regressor Results:
Mean Squared Error (MSE): 3586350.75
Mean Absolute Error (MAE): 1427.26
Median Absolute Error: 1079.79
 R^2 Score: 0.77
Explained Variance Score: 0.77

